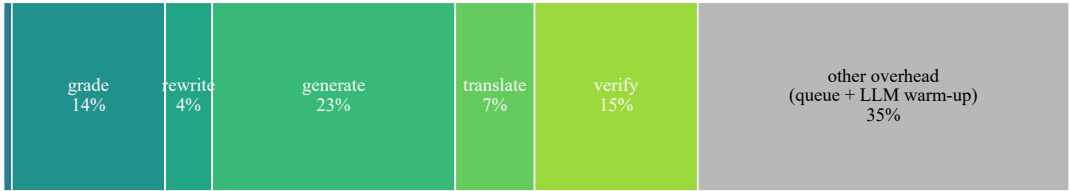
grade
14%
rewrite
4%
generate
23%
translate
7%
verify
15%
other overhead
(queue + LLM warm-up)
35%
0
10000
20000
30000
40000
50000
End-to-end latency (ms) — mean = 55.4 s · p50 = 53.0 s · p95 = 83.6 s (CPU-only LLM inference)