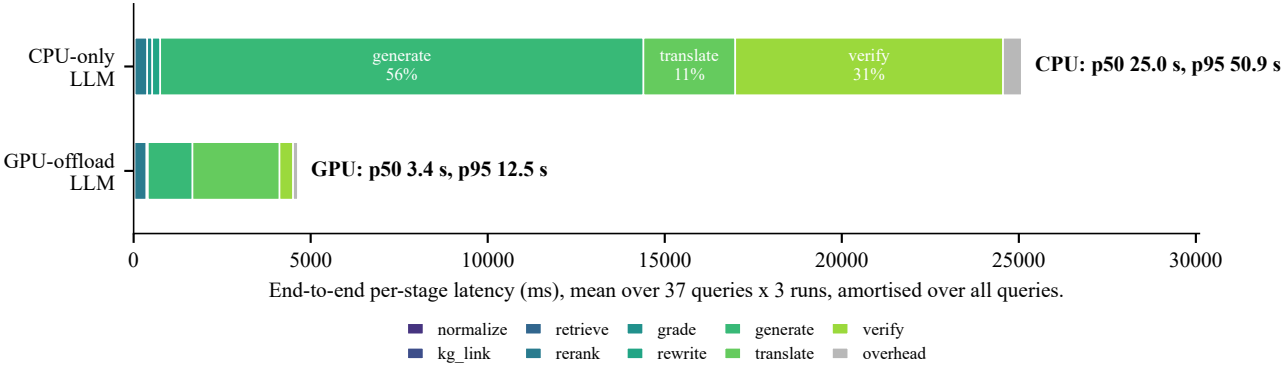
CPU-only LLM
generate 56%
translate 11%
verify 31%
CPU: p50 25.0 s, p95 50.9 s
GPU-offload LLM
GPU: p50 3.4 s, p95 12.5 s
0
5000
10000
15000
20000
25000
30000
End-to-end per-stage latency (ms), mean over 37 queries x 3 runs, amortised over all queries.
normalize
kg_link
retrieve
rerank
grade
rewrite
generate
translate
verify
overhead